

End-to-End GPU Ingestion: Tokenization to VRAM

